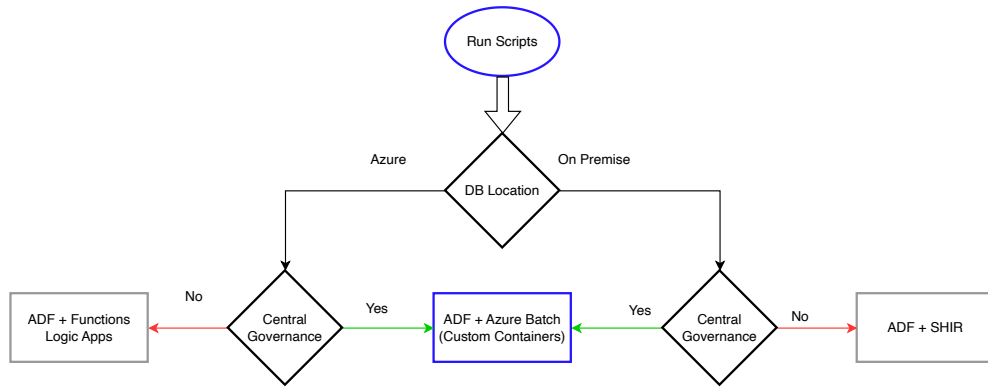


❌ 1. Azure Data Factory does **not natively support executing arbitrary scripts** — including SQL, Bash, Python, or PowerShell — on PostgreSQL targets, unlike its native support for Azure SQL via stored procedure and script activities.

⚠️ 2. While the PostgreSQL Connector v2.0 supports authentication via **Service Principal**, its functionality is **limited to data movement operations** such as **CopyActivity** and **metadata retrieval** (LookupActivity); it does not support the execution of custom DDL/DML SQL scripts.

✅ 3. Future versions of the PostgreSQL connector (e.g., v3.0 and beyond) **may extend functionality** to support custom script execution and additional activities beyond data movement; **however, as of now, such capabilities are not available.**



#### ⚠️ 1. Limited Script Execution Scope and Flexibility

- Azure Functions and Logic Apps are well-suited for **lightweight, event-driven tasks** but are **not ideal for full SQL script execution**, especially DDL-heavy workloads.
- You'll need to manually handle psql or custom connectors. — they **don't natively support arbitrary script execution** against PostgreSQL like Batch can.

#### ⚠️ 2. Complex Authentication and Execution Context

- Managed Identity support exists but **requires explicit binding per app/function**, with more complexity in **multi-tenant or hybrid setups**.
- **Execution context is stateless and time-bound**, which can be problematic for long-running or dependent script chains.

#### ⚠️ 3. Limited Control, Observability, and Reusability

- Functions and Logic Apps offer **limited control over runtime environments** — you can't containerize or bundle custom tooling.
- Debugging, logging, and retry handling are **less granular** than Azure Batch, and scaling beyond basic scenarios often introduces more orchestration overhead.

#### ✅ 1. Centralized, Scalable, and Consistent Execution

- Enables **centralized job execution** across all regions, tenants, and network zones using reusable Batch pools.
- Custom containers provide a **consistent runtime** (e.g., with psql, Python, tools pre-installed), eliminating environment drift.

#### ✅ 2. Strong Governance and Secure Identity Management

- Supports **Service Principal or Managed Identity** authentication at the pool level — easy to audit, rotate, and control via Azure RBAC.
- Batch jobs are executed in **isolated, policy-governed containers**, enabling centralized control and logging.

#### ✅ 3. Flexible Networking and Tooling

- Easily integrates with **VNETs and private endpoints**, suitable for both cloud and hybrid (on-prem) PostgreSQL targets.
- Custom containers allow you to **bundle all required scripts, binaries, and tools**, reducing operational dependencies and setup time.

#### ❌ 1. Decentralized Execution and Governance

- SHIR requires per-environment installation, making it difficult to enforce **centralized control**, policy management, or consistent access via service principals.
- No centralized execution pool or unified authentication strategy across tenants, regions, or on-prem environments.

#### ❌ 2. Operational Complexity and Poor Scalability

- High setup and maintenance overhead: each SHIR needs **manual installation**, patching, and capacity planning.
- **Does not scale well** for large estates — every new DB or network zone may require a separate SHIR or config.

#### ❌ 3. Rigid Networking and No Runtime Flexibility

- Tightly coupled to the network layer — **must reside within the same VNET or firewall scope** as the target DB.
- Cannot run custom containers or pre-built execution environments like Azure Batch, making **tooling and dependency management** inflexible and error-prone.